

# DISTRIBUTED COMPUTING IN SENSOR NETWORKS


James Long<sup>1</sup> Anupa Murali<sup>2</sup> Zihao Wang<sup>3</sup>

<sup>1</sup>PhD Candidate  
Civil Engineering  
MIT

<sup>2</sup>Affiliation  
Harvard University

<sup>2</sup>Affiliation  
Harvard University

# BACKGROUND



- Mode shapes describe the dynamic response of mechanical systems to their environment
- Material properties (mass, stiffness, damping), can be inferred from these responses
- Changes in modal responses may indicate defects

# COMPUTATION

Given time series  $x(t)$  from  $n$  sensor locations:

Compute  $X_n(\omega)$  for each of the  $n$  locations

For a specific  $\omega_i$  form  $G_{n \times n}(\omega_i)$  where:


$$G_{jk}(\omega_i) = X_j(\omega_i) \times X_k(\omega_i)^*$$

Then decompose  $G_{n \times n} = U_{n \times n} \Sigma_{n \times n} V_{n \times n}^*$

The first singular vector (column) of  $U$  is the mode shape estimate

# MOTIVATION: TECHNOLOGY CAN OVERWHELM TRADITIONAL ALGORITHMS

Example 1: high-res cameras can produce huge data sets  
(e.g one time series per pixel)



# MOTIVATION: TECHNOLOGY CAN OVERWHELM TRADITIONAL ALGORITHMS

Example 2: Networks of wirelessly connected sensor devices

- Devices have built in accelerometers and microcontrollers
- Very low cost enables large scale, dense deployment
- Communication out of the network drains battery, and we are forced to perform computation on slow (100MHz), memory constrained (128kB RAM) devices

# APPROACH & SECRET WEAPON

- Recall that we use only the first singular vector as an estimate of the mode shape
- Because our matrix is normal, we can use a power method to estimate the largest eigenvector (PageRank!)

Algorithm:

$$i = 0; \text{ init } v_0$$

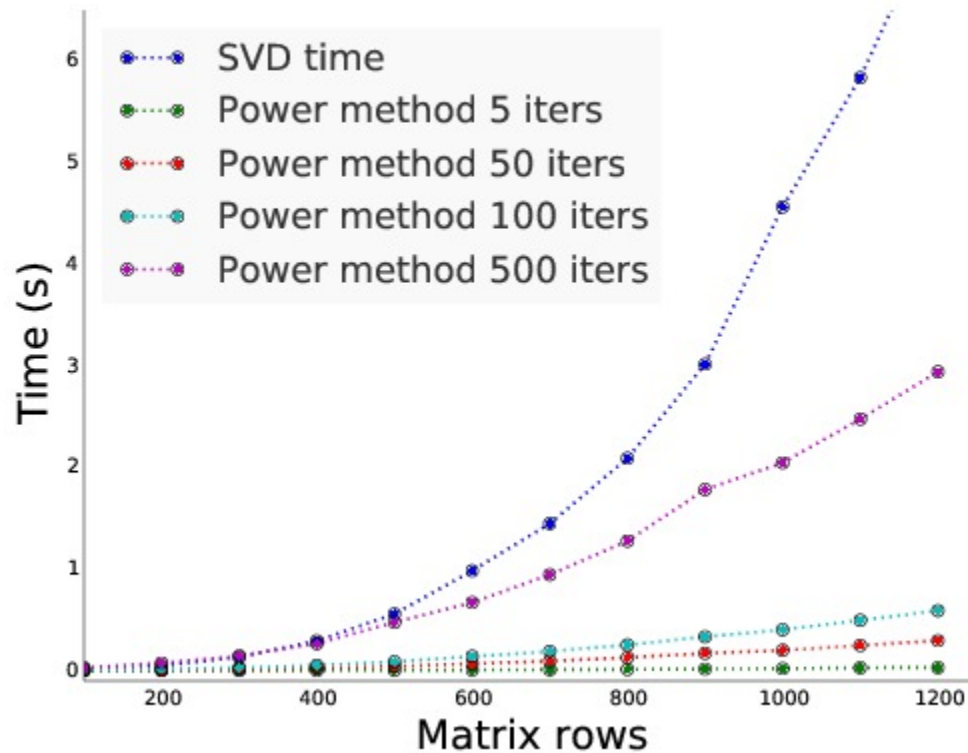
while ( $i < N$ ) :

$$v_{i+1} = Gv_i$$

$$v_{i+1} = \frac{v_{i+1}}{\|v_{i+1}\|}$$

$$i = i + 1$$

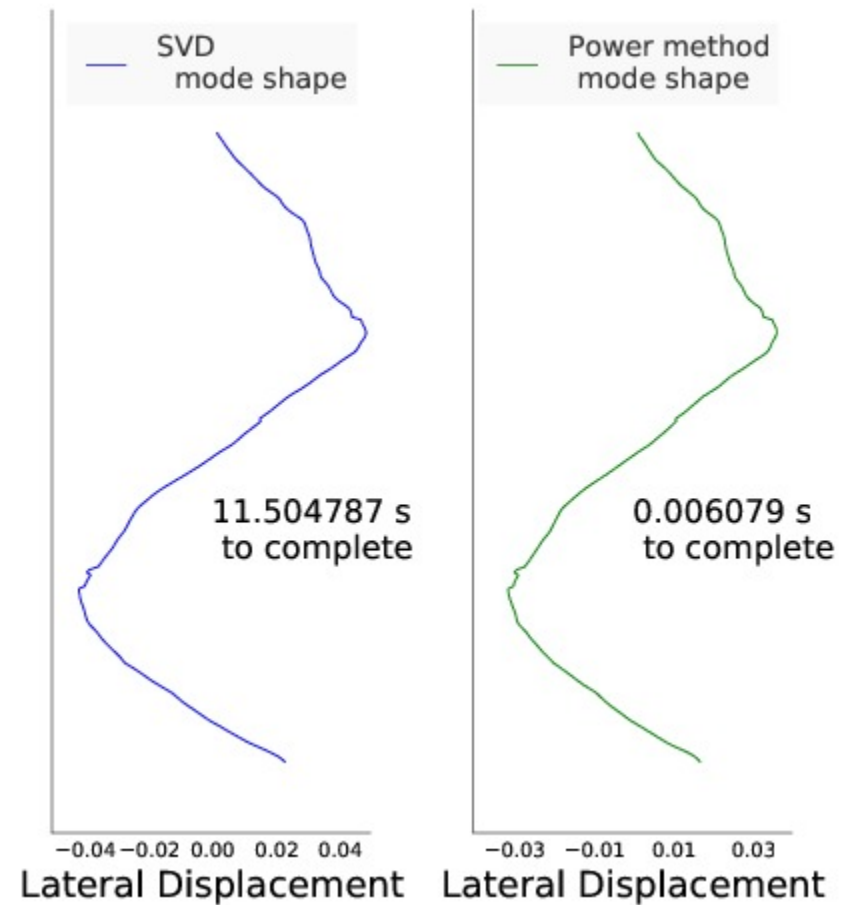

# APPROACH & SECRET WEAPON



Rate of convergence:  $\frac{\lambda_1}{\lambda_2}$

$\lambda_1$  is the largest eigenvalue,  $\lambda_2$  the next largest

# RESULTS: SERIAL COMPARISON




1217 individual time series of  
length 1000

1900x speed increase

$$\lambda_1 = 3.9 \times 10^5 \quad \lambda_2 = 1.2 \times 10^{-9}$$



# RESULTS: SERIAL COMPARISON

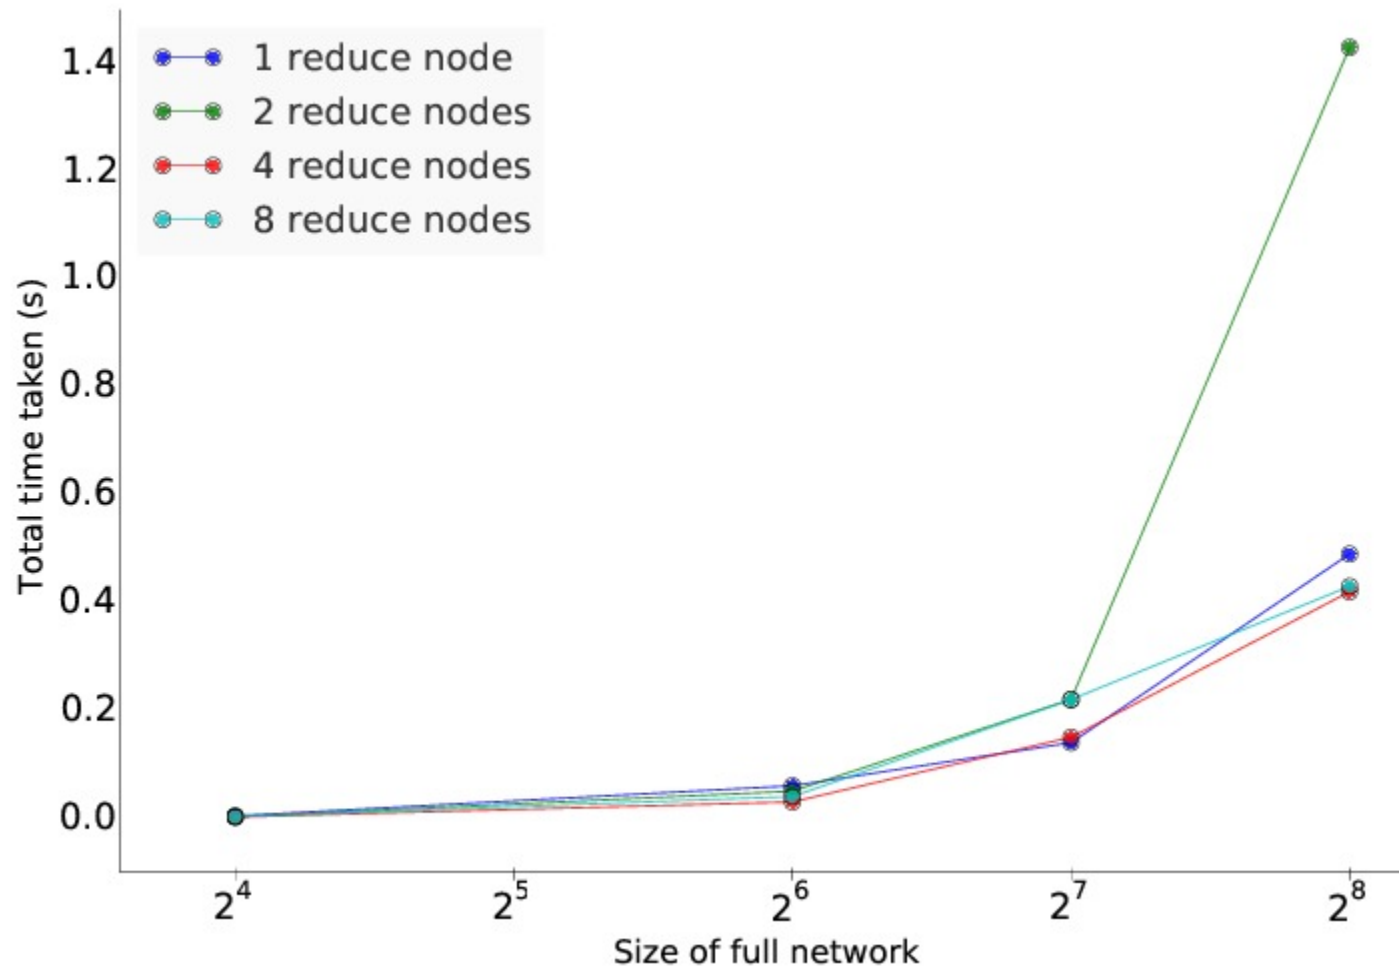


54 individual time series of  
length 180000

# RESULTS: PARALLELISATION IN MPI

- MPI offers a solution to very large data sets, as well as a good simulation of a sensor network (using send/recv)
- We can exploit embarrassing parallelism of FFT calculation: One process for every time series
- Then collect FFT values at a subset of processes ('reduce'), otherwise communication will kill us
- Question: What is the right number of 'reduce' nodes

# RESULTS: MPI ODYSSEY BENCHMARKING



Note: A realistic serial implementation would be orders of magnitude more costly

# EVALUATION

- Using the power method to estimate mode shapes, instead of SVD saves orders of magnitudes in execution time on CPU ( $1900\times$  on  $1217\times 1217$  matrix)
- No meaningful decrease in accuracy
- Successfully parallelised the algorithm using MPI4py: Simulation of a wireless network
- More work required to prove concept with wireless communication and microcontroller

# MAIN LESSON

- Parallelism can help immensely - or hurt
- We need to decide prudently how much parallelism to use (MapReduce is a good example)
  - Arithmetic intensity is important!
    - Good algorithms are crucial